

Modern CI/CD for Static Websites: Automating Deployment with Docker, Nginx, and Shell Scripts

All rights reserved.

This repository and its contents are private and intended solely for users authorized by ISWAD Inc. Unauthorized access, use, modification, or distribution is strictly prohibited.

For more information, visit the repository at: <https://github.com/mmmohajer/ci-cd-static-sites>

Table of Contents

- [Purpose](#)
- [How to Use This Repository and Study](#)
- [Who Is This For?](#)
- [What You'll Learn](#)
- [Prerequisites](#)
- [Step 01: Preparing the Development Configuration](#)
 - [Concepts to Learn](#)
 - [Practical Steps](#)
 - [Create the Folder Structure](#)
 - [Write the Nginx Configuration](#)
 - [Write the Docker Compose File](#)
 - [Folder Structure](#)
 - [Run the App](#)
 - [Verify Everything is Working](#)
 - [Test Your Setup](#)
- [Step 02: Serving Static Files and Improving Performance with Gzip](#)
 - [Concepts to Learn](#)
 - [Practical Steps](#)
 - [Add Static File Support in Nginx](#)
 - [Add Gzip Configuration](#)
 - [Update Your Dockerfile.dev](#)
 - [Add Static Files](#)
 - [Folder Structure](#)
 - [Run and Test the Setup](#)
 - [Test Your Setup](#)
- [Step 03: Setting Up the Production Environment](#)
 - [Goal of This Step](#)
 - [Prerequisites](#)
 - [Steps to Configure the Production Environment](#)
 - [Create the Nginx Configuration for SSL Creation](#)
 - [Create a Dockerfile for SSL Setup](#)
 - [Create the Docker Compose File](#)
 - [Create the init-letsencrypt.sample.sh Script](#)
 - [SSH into the Server](#)

- Set Up Git Configuration
 - Deploy the App and Create SSL
 - Verify SSL Creation
 - Folder and File Structure
 - Step 04: Running Application on The Server
 - Goal of This Step
 - Steps to Configure the Production Environment
 - Create the Nginx Configuration for Production
 - Create the Dockerfile for Production
 - Create the Docker Compose File for Production
 - Create the Deployment Script
 - Push the Changes to GitHub
 - Deploy the App on the Server
 - Verify Your App
 - Automate Certificate Renewal
 - Folder and File Structure
 - Conclusion
-

Purpose

This project is a step-by-step tutorial designed to simulate how CI/CD pipelines are built and scaled for real-world applications. Starting from the basics, you'll learn how to set up a robust CI/CD pipeline for static websites using Docker, Nginx, and Shell scripts. By following each step, you'll understand how to scale from a simple setup to a fully functional CI/CD pipeline ready for production.

This setup is reusable for any static website, whether it's:

- A personal portfolio website.
- A website to represent your services.
- A business website needing regular updates.

By the end of this tutorial, you'll have a foundational CI/CD pipeline that shifts your focus from deployment headaches to content creation and design updates. You'll no longer worry about how to:

- Serve your website.
- Set up a custom domain.
- Integrate SSL certificates.

Instead, your focus will shift entirely to creating and updating the content and beautifying your website with HTML, CSS, and JavaScript.

How to Use This Repository and Study

This tutorial is designed as a **step-by-step guide**, with each step building upon the previous one. To get the most out of this project, follow these instructions:

1. Follow Each Step Sequentially

- Each step is carefully structured to simulate the process of building and scaling a CI/CD pipeline in real-world applications.
- **Do not skip steps**, as every new step builds on the work from the previous one.

2. Instructions for Each Step

- In each step, the **README.md** will provide:
 - **Tasks to Complete:** Specific actions you need to perform (e.g., creating files, writing code).
 - **Tests:** A way to verify that you have completed the step correctly.

3. Encourage Self-Learning

- If you encounter challenges during a step:
 - **Google it:** Use search engines to find solutions.
 - **Use ChatGPT:** Ask for guidance or clarification.
 - **Visit Stack Overflow:** Research common issues and solutions.
- Spending time troubleshooting on your own is an invaluable part of the learning process.

4. Get Support When Needed

If you're still unable to resolve an issue after self-search:

- **Contact Us:** Send a support request through our app for assistance.
- **Book a Mentorship Session:** For more in-depth help, you can schedule a mentorship session with us via the app.
- For more information about these services, feel free to reach us at **info@iswad.tech**.

5. Use the Branches for Guidance

- Each step has a corresponding branch named **step-{STEP_NUMBER}** (e.g., **step-01**, **step-02**):
 - Switch to the relevant branch to see how the step should be structured:

```
git checkout step-{STEP_NUMBER}
```

- Review the comments in the code, where every line or important section is explained.
- If you want to compare your solution with mine, you can pull the code from the branch:

```
git pull origin step-{STEP_NUMBER}
```

6. Compare Changes Between Steps

- To fully understand what changed and why, use Git to compare files and folders updated in the current step:
 - **Using Pull Requests:** If you're using a repository with pull requests for each step, review the diff to analyze updates.
 - **Using Command Line:**

- Compare changes between steps:

```
git diff step-{PREVIOUS_STEP_NUMBER} step-  
{CURRENT_STEP_NUMBER}
```

- This will show the differences between the two steps, helping you understand the updates made.

7. Track Your Progress

After completing each step:

- **Mark the Step as Completed:** Update the status of that step in the app to indicate that it has been successfully completed.
- **Share Optional Feedback:**
 - Include what you learned during the step or what challenges you faced and how you solved them.
 - This information will help:
 - Keep track of your progress and make you more committed to your work.
 - Assess whether you have successfully completed the step.
 - Enable future AI-powered features (currently under development) that will:
 - Help you generate a better resume.
 - Suggest relevant topics and personalized roadmaps to achieve your goal of becoming a professional or senior developer.

By following this structured process, you'll gain both practical knowledge and confidence in building CI/CD pipelines for static websites. Additionally, our tools and features will help you stay motivated, track progress, and work towards your long-term career goals.

Who Is This For?

This project is for anyone who wants to learn how to automate and scale website deployments, including:

- **DevOps Engineers:** Build expertise in automating pipelines for static websites.
 - **Full-Stack Engineers:** Level up by learning the basics of DevOps.
 - **Front-End Developers:** Learn how to publish your designs professionally.
 - **UI/UX Designers:** If you can export HTML, CSS, and JavaScript or have basic knowledge of it, this guide will show you how to publish your app.
 - **Professionals and Service Providers:** Build and scale your website for your services or business.
-

What You'll Learn

By completing this tutorial, you'll gain:

- The ability to create and manage **Dockerfiles** for static websites.

- Hands-on experience configuring **Nginx** for efficient static file serving.
 - Skills to write **Shell scripts** for automating CI/CD pipelines.
 - An understanding of CI/CD concepts and how to apply them to real-world projects.
 - A reusable deployment pipeline that works for any static website.
-

Prerequisites

Before you begin, ensure you have the following installed:

- **Git:** [Install Git](#).
 - **GitHub account:** For managing the repository and hosting code.
 - **Docker:** [Install Docker](#).
 - **Shell:** A Unix-based shell (bash/zsh) for running scripts.
 - **Basic Knowledge:** Familiarity with HTML, CSS, and JavaScript.
-

Step 01: Preparing the Development Configuration

Goal of This Step

In this step, you will:

1. Set up the development environment for your static website.
2. Learn the basics of **Nginx**, **Docker**, and **Docker Compose** to containerize and serve the app.
3. Run the app locally using `docker-compose` and verify that everything is working as expected.

By the end of this step, you'll have a minimal CI/CD setup running locally, serving your static website.

Concepts to Learn

1. Nginx: The Backbone of Modern Web Servers

Theoretical Overview: Nginx is a web server used to efficiently serve files (like HTML, CSS, and JavaScript) or act as a reverse proxy. It's widely used in real-world scenarios for its performance and ability to handle thousands of simultaneous requests.

Practical Analogy: Think of Nginx as a **waiter** in a restaurant:

- When a user (customer) makes a request, like "Show me the homepage," Nginx fetches the `index.html` file (menu) and delivers it to the user.
- If the customer requests additional resources (like CSS or images), Nginx fetches and serves them quickly.

Real-World Example:

- Websites like **Netflix**, **YouTube**, and **Amazon** use Nginx to serve static files like videos, images, and stylesheets.

How Nginx Works in This Project:

- It will:
 1. Serve your `index.html` file to users visiting the site.
 2. Look for additional static files (like CSS and JavaScript) and deliver them to the browser.
-

2. Docker: Ensuring Consistency

Theoretical Overview: Docker packages your app and its environment into a container, ensuring it runs the same way across all machines, from your laptop to a production server.

Practical Analogy: Think of Docker as a **portable kitchen** that includes:

- The ingredients (your app files).
- The tools (dependencies like Nginx).
- The oven (a runtime environment).

Real-World Example:

- Companies like **Spotify** and **Google** use Docker to make sure their apps run identically across development and production environments.

How Docker Works in This Project:

- It will:
 1. Package your app files and Nginx into a container.
 2. Serve your app consistently, no matter where it's running.
-

3. Docker Compose: Managing Multiple Services

Theoretical Overview: Docker Compose lets you define and manage multiple services in one file (`docker-compose.yml`) and run them together with a single command.

Practical Analogy: Think of Docker Compose as a **kitchen manager** that coordinates multiple stations:

- One station serves files (Nginx).
- Another could handle data storage (a database).

Real-World Example:

- An e-commerce platform might use Docker Compose to manage:
 - A front-end service (React or Angular).
 - A back-end API (Node.js or Django).
 - A database (PostgreSQL or MongoDB).

How Docker Compose Works in This Project:

- It will:
 1. Define the Nginx service.
 2. Map your local files to the container for live updates.
 3. Expose port `80` so you can access the app in your browser.
-

Practical Steps

1. Create the Folder Structure

- At the root of your repository:
 - Create a folder named `site` to store all your website files (HTML, CSS, JavaScript, images).
 - Add a file named `index.html` with minimal content (e.g., "Hello, world!").
 - Add a file named `404.html` to serve as a custom 404 error page.
 - Create a folder named `nginx` to store Nginx configuration files.
 - Add a file named `default-dev.conf` for your Nginx configuration.
 - Add a file named `Dockerfile.dev` for building the Nginx Docker image.
 - Create a file named `docker-compose-dev.yml` in the root directory for your Docker Compose setup.
-

2. Write the Nginx Configuration

- Inside the `nginx/default-dev.conf` file, write a basic configuration to:
 - Serve files from `/var/www/app`.
 - Look for `index.html` as the default file when users visit the root of your site.
 - Include a fallback to a `404.html` file for missing resources.

```
server {  
    # The server listens for incoming HTTP requests on port 80 (default  
    for HTTP).  
    listen 80;  
  
    # Defines the root directory where the website files are stored and  
    served from.  
    root /var/www/app;  
  
    # Specifies the default file to serve when a user visits the root URL.  
    index index.html;  
  
    # Defines a custom 404 error page to display when a file is not found.  
    error_page 404 /404.html;  
  
    # Handles all other requests to the root or subpaths.  
    location / {  
        # Attempts to serve the requested file or path in this order:  
        # 1. An exact file match (e.g., /about.html).  
        # 2. A file with .html appended (e.g., /about -> /about.html).  
        # 3. A directory with the name (e.g., /about/).  
        # 4. If none of the above exist, serves the custom 404 error page.  
        try_files $uri $uri.html $uri/ /404.html;  
    }  
}
```

For building nginx image with the appropriate configuration, inside Dockerfile.dev file add:

```
# This Dockerfile sets up an Nginx container for serving static files in a
development environment.
# - It uses a lightweight Alpine-based Nginx image.
# - Copies a custom Nginx configuration to the container.
# - Prepares a directory for static files with appropriate permissions.

# Use a lightweight Nginx image based on Alpine Linux for efficiency.
FROM nginx:1.20.2-alpine

# Copy the custom Nginx configuration file # to the container's default
configuration location.
COPY ./default-dev.conf /etc/nginx/conf.d/default.conf

# Create a directory for static files within the container.
# Set permissions to allow read, write, and execute for the owner, and
read/execute for others.
RUN mkdir -p /var/www/app/static && \
    chmod 755 /var/www/app/static
```

3. Write the Docker Compose File

- In the root directory, define the `docker-compose-dev.yml` file to:
 - Set up Nginx as a service.
 - Map your local `site` folder to `/var/www/app` inside the container.
 - Map the `nginx` folder for custom configurations.
 - Expose port `80` for local access.

```
# This Docker Compose file defines a service for running an Nginx
container in development mode.
# It builds a custom Docker image for Nginx using a Dockerfile located in
the "nginx" folder.
# The "site" folder on the host is mapped into the container to serve
static website files.
# Port 80 on the host is mapped to port 80 in the container, allowing
access via http://localhost.
```

```
services:
  nginx: # Defines a service named "nginx"
    restart: always # Ensures the container restarts automatically if it
stops or crashes
    build: # Configuration for building the Docker image
      context: ./nginx # Specifies the build context (folder containing
Dockerfile and related files)
      dockerfile: Dockerfile.dev # Specifies the Dockerfile to use for
building the image
    ports:
```



```
- "80:80" # Maps port 80 on the host to port 80 in the container
(access via http://localhost)
volumes:
- ./site:/var/www/app # Maps the "site" folder on the host to
"/var/www/app" inside the container
```

4. Folder Structure

At the end of this step, your folder structure should look like this:

```
.
├── nginx/
│   ├── default-dev.conf # Write your Nginx configuration in this file
│   └── Dockerfile.dev   # Dockerfile to build the Nginx image
├── site/
│   ├── index.html       # Write your minimal HTML file here (e.g.,
│   │                   "Hello, world!" app)
│   └── 404.html         # Custom 404 error page for missing resources
└── docker-compose-dev.yml # Define your Docker Compose setup for
development
```

5. Run the App

- Open a terminal or Bash at the root directory where `docker-compose-dev.yml` is located.
- Run the following command to start the app:

```
docker-compose -f docker-compose-dev.yml up --build -d
```

6. Verify Everything is Working

1. Visit <http://localhost> in your browser. You should see the content of your `index.html` file.
2. Modify the `index.html` file to display a "Hello, world!" message.
3. Refresh the browser to ensure the changes are reflected.

6. Test Your Setup

After completing this step, verify the following:

- Does visiting <http://localhost> display the content of your `index.html` file?
 - Are changes to the `index.html` file reflected in the browser after refreshing?
-

Step 02: Serving Static Files and Improving Performance with Gzip

Goal of This Step

In this step, you will:

1. Update your Nginx configuration to serve static files (e.g., CSS, JS, images) from a dedicated `static` folder.
2. Optimize your website's performance by enabling Gzip compression.
3. Test the setup by adding CSS and JavaScript files to the `static` folder and verifying that they are served correctly.

By the end of this step, you'll have a more functional setup capable of serving all your static assets efficiently with improved performance.

Concepts to Learn

1. Static File Serving

Nginx can efficiently serve static files like CSS, JS, and images directly from a specific folder. This avoids unnecessary processing and delivers these files to the user quickly.

- **Practical Example:**
 - Websites like **Amazon** and **Netflix** use Nginx to serve static assets (e.g., stylesheets, images, and JavaScript) from dedicated directories.
- **How It Works Here:**
 - We'll configure Nginx to map `/static/` URLs to the `static` folder inside your project directory (`/var/www/app/static`).

2. Gzip Compression

Gzip reduces the size of text-based files (like HTML, CSS, and JavaScript) before sending them to the browser, resulting in faster load times.

- **Practical Example:**
 - Gzip is widely used by high-traffic sites to save bandwidth and improve performance.
 - **How It Works Here:**
 - We'll enable Gzip in Nginx and configure it to compress supported file types like `text/css`, `application/javascript`, and `text/html`.
-

Practical Steps

1. Add Static File Support in Nginx

1. Update your `nginx/default-dev.conf` file to include the following blocks:

```
location /static/ {
    # Maps requests for /static/ to the static folder in your
    project.
    alias /var/www/app/static;
}

location ~* \.(gif|jpe?g|png|svg)$ {
    # Caches these file types for as long as possible.
    expires max;
    # Ensures long-term caching.
    add_header Cache-Control "public, max-age=31536000, immutable";
}

location ~* \.(css|js|json)$ {
    # Ensures proper caching behavior.
    add_header Cache-Control "no-cache, must-revalidate";
}
```

2. Add Gzip Configuration

- Create a new file in the `nginx` folder named `gzip.conf`.
- Add the following lines to enable Gzip:

```
gzip on;
gzip_types text/plain text/css application/json application/javascript
text/xml application/xml application/xml+rss text/javascript;
gzip_vary on;
gzip_proxied any;
# Compression level: balance between performance and CPU usage.
gzip_comp_level 6;
gzip_buffers 16 8k;
gzip_http_version 1.1;
# Only compress files larger than 256 bytes.
gzip_min_length 256;
```

Update Your Dockerfile.dev to Include the Gzip Configuration

Add the following line to your `Dockerfile.dev`:

```
# Copies the Gzip configuration into the container.
COPY ./gzip.conf /etc/nginx/conf.d/gzip.conf
```

3. Add Static Files

1. Inside the `site` folder, create a `static` directory.

2. Inside the static folder, create two subfolders css, js
3. Add some files to test:
 - Inside css folder, create a file named `styles.css` with the following content:

```
body {  
  background-color: blue;  
  font-family: Arial, sans-serif;  
  color: wheat;  
}
```

- Inside js folder, create a file named `main.js` with the following content:

```
console.log("Hello, world!");
```

4. Link these files in your `index.html`:

```
<!DOCTYPE html>  
<html lang="en">  
  <head>  
    <meta charset="UTF-8" />  
    <meta name="viewport" content="width=device-width, initial-  
scale=1.0" />  
    <title>Hello World!</title>  
    <link rel="stylesheet" href="/static/css/styles.css" />  
  </head>  
  <body>  
    Hello World!!!  
  
    <script src="/static/js/main.js"></script>  
  </body>  
</html>
```

4. Folder Structure

After completing this step, your folder structure should look like this:

```
.  
├── nginx/  
│   ├── default-dev.conf # Updated Nginx configuration for static files  
│   ├── gzip.conf       # Gzip compression configuration  
│   └── Dockerfile.dev   # Dockerfile for the Nginx service  
├── site/  
│   ├── index.html      # Updated HTML linking static files  
│   └── static/
```

```
├── css/
│   └── styles.css # Test CSS file
├── js/
│   └── main.js    # Test JavaScript file
└── docker-compose-dev.yml # Docker Compose file for development
```

5. Run and Test the Setup

1. Restart your Docker containers to apply the changes:

```
docker-compose -f docker-compose-dev.yml up --build -d
```

Open your browser and visit:

- <http://localhost>: You should see your updated `index.html` with the applied CSS styles.
- Check the browser's developer tools (**Console tab**) to confirm that the JavaScript file is executed.

Verify Gzip compression:

- Use an online tool like [Gzip Checker](#) or your browser's **Network tab** to confirm that files like `styles.css` and `scripts.js` are compressed.

6. Test Your Setup

- Does visiting <http://localhost> show the updated page with CSS and JavaScript applied?
- Are static files served from the `/static/` folder?
- Is Gzip compression applied to your text-based files (CSS, JS, HTML)?

By completing this step, you've set up a static file server with Nginx and optimized it with Gzip compression for faster delivery. This serves as the foundation for efficient static asset management in your CI/CD pipeline.

Step 03: Setting Up the Production Environment

Goal of This Step

In this step, we will:

1. Set up a production-ready environment with a real domain.
2. Configure a Linux Ubuntu server to host the application.
3. Generate SSL certificates for the domain using Certbot to enable HTTPS.
4. Build and deploy a minimal Nginx-based setup for SSL creation.

By the end of this step, you'll have the foundation ready to run your app on a production server with a valid SSL certificate.

Prerequisites

Before proceeding:

1. **Purchase a Domain:**
 - Use providers like **GoDaddy** or others to buy a domain.
 2. **Set Up a Linux Server:**
 - Purchase a Linux server (preferably Ubuntu) from a provider like **DigitalOcean**.
 - DigitalOcean is recommended for its simplicity, features, and pricing.
 3. **DNS Configuration:**
 - Create an **A record** in your domain's DNS settings pointing to your server's public IP address.
-

Steps to Configure the Production Environment

1. Create the Nginx Configuration for SSL Creation

1. In the `nginx` folder, create a file named `default-create-ssl.conf` with the following content:

```
server {
    # Configures the server to listen on port 80 for HTTP requests.
    listen 80;

    # Replace `APP_URL` with your actual domain name (e.g., example.com).
    # This specifies the domain that this Nginx configuration will serve.
    server_name APP_URL www.APP_URL;

    # Disables the Nginx version number in error pages and headers for
    security.
    server_tokens off;

    location /.well-known/acme-challenge/ {
        # Allows unrestricted access to the `.well-known/acme-challenge`
        directory,
        # which is required for Certbot to verify domain ownership.
        allow all;

        # Specifies the directory where Certbot will store temporary
        challenge files
        # used to verify the domain when generating an SSL certificate.
        root /var/www/certbot;
    }

    location / {
        # Serves static HTML files from the `/usr/share/nginx/html`
        directory.
        root /usr/share/nginx/html;

        # Specifies the default files (index.html or index.htm) to serve
        # when a user visits the root URL.
        index index.html index.htm;
    }
}
```

```
# Attempts to serve the requested file (`$uri`).  
# If the file is not found, it falls back to `/index.html`.  
try_files $uri $uri/ /index.html;  
}  
}
```

2. Create a Dockerfile for SSL Setup

1. Navigate to the `nginx` folder in your project directory.
2. Create a file named `Dockerfile.create-ssl` and add the following content:

```
# Use the lightweight Alpine-based Nginx image as the base image.  
FROM nginx:1.20.2-alpine  
  
# Copy the minimal Nginx configuration for SSL creation into the  
container.  
COPY ./default-create-ssl.conf /etc/nginx/conf.d/default.conf
```

3. Create the Docker Compose File

1. Navigate to the root folder of your project.
2. Create a file named `docker-compose-create-ssl.yml` and add the following content:

```
services:  
  nginx:  
    restart: always  
    # Ensures the Nginx container restarts automatically if it stops  
    or crashes.  
  
    build:  
      context: ./nginx  
      # Specifies the build context (the directory containing the  
      Dockerfile).  
      dockerfile: Dockerfile.create-ssl  
      # Defines the Dockerfile to use for building the Nginx image.  
  
    ports:  
      - "80:80"  
      # Maps port 80 on the server to port 80 in the container,  
      allowing HTTP traffic.  
  
    volumes:  
      - ./site:/usr/share/nginx/html  
      # Maps the local `site` folder to `/usr/share/nginx/html` in the  
      container.  
      # This folder contains the static files (e.g., index.html) for
```

the Nginx server.

```
- ./nginx/certbot/www:/var/www/certbot
# Maps the local folder for Certbot's `.well-known` directory to
`/var/www/certbot` in the container.
# Certbot uses this directory to store temporary challenge files
for domain verification.
```

```
- ./nginx/certbot/conf:/etc/letsencrypt
# Maps the local folder for Certbot's configuration and SSL
certificates to `/etc/letsencrypt` in the container.
# This is where Certbot stores the generated SSL certificates.
```

```
certbot:
  image: certbot/certbot:latest
  # Uses the latest Certbot image from Docker Hub to manage SSL
  certificates.
```

```
  container_name: certbot
  # Sets the name of the Certbot container for easier
  identification.
```

```
  volumes:
    - ./nginx/certbot/www:/var/www/certbot
    # Shares the `.well-known` folder with the Certbot container for
    domain verification.
```

```
    - ./nginx/certbot/conf:/etc/letsencrypt
    # Shares the configuration and SSL certificate folder with the
    Certbot container.
```

```
  entrypoint: "/bin/sh -c 'trap exit TERM; while ;; do sleep 6h &
  wait ${!}; certbot renew; done'"
  # Defines the custom entrypoint for the Certbot container:
  # - Keeps the container running in a loop.
  # - Attempts to renew SSL certificates every 6 hours.
  # - Ensures the process exits gracefully if the container stops.
```

4. Create the `init-letsencrypt.sample.sh` Script

1. In the root folder of your project, create a file named `init-letsencrypt.sample.sh`.
2. Make the script executable:

```
chmod +x init-letsencrypt.sh
```

3. Add the following code to the file:


```
#!/bin/bash

if ! [ -x "$(command -v docker-compose)" ]; then
echo 'Error: docker-compose is not installed.' >&2
exit 1
fi

# Replace APP_URL with your actual domain (e.g., example.com)
domains=(APP_URL www.APP_URL)

rsa_key_size=4096
data_path="./nginx/certbot"

# Update EMAIL for Let's Encrypt notifications
email="EMAIL"

staging=0

if [ ! -e "$data_path/conf/options-ssl-nginx.conf" ] || [ ! -e
"$data_path/conf/ssl-dhparams.pem" ]; then
echo "### Downloading recommended TLS parameters ..."
mkdir -p "$data_path/conf"
curl -s https://raw.githubusercontent.com/certbot/certbot/master/certbot-
nginx/certbot_nginx/_internal/tls_configs/options-ssl-nginx.conf >
"$data_path/conf/options-ssl-nginx.conf"
curl -s
https://raw.githubusercontent.com/certbot/certbot/master/certbot/certbot/s
sl-dhparams.pem > "$data_path/conf/ssl-dhparams.pem"
echo
fi

echo "### Creating dummy certificate for $domains ..."
path="/etc/letsencrypt/live/$domains"
mkdir -p "$data_path/conf/live/$domains"
docker-compose -f ./docker-compose-create-ssl.yml run --rm --entrypoint "\
openssl req -x509 -nodes -newkey rsa:$rsa_key_size -days 1\
-keyout '$path/privkey.pem' \
-out '$path/fullchain.pem' \
-subj '/CN=localhost'" certbot
echo

echo "### Starting nginx ..."
docker-compose -f ./docker-compose-create-ssl.yml up --force-recreate -d
nginx
echo

echo "### Deleting dummy certificate for $domains ..."
docker-compose -f ./docker-compose-create-ssl.yml run --rm --entrypoint "\
rm -Rf /etc/letsencrypt/live/$domains && \
rm -Rf /etc/letsencrypt/archive/$domains && \
rm -Rf /etc/letsencrypt/renewal/$domains.conf" certbot
```

```

echo

echo "### Requesting Let's Encrypt certificate for $domains ..."
#Join $domains to -d args
domain_args=""
for domain in "${domains[@]}; do
domain_args="$domain_args -d $domain"
done

# Select appropriate email arg
case "$email" in
"") email_arg="--register-unsafely-without-email" ;;
*) email_arg="--email $email" ;;
esac

# Enable staging mode if needed
if [ $staging != "0" ]; then staging_arg="--staging"; fi

docker-compose -f ./docker-compose-create-ssl.yml run --rm --entrypoint "\
certbot certonly --webroot -w /var/www/certbot \
    $staging_arg \
    $email_arg \
    $domain_args \
    --rsa-key-size $rsa_key_size \
    --agree-tos \
    --force-renewal" certbot
echo

echo "### Reloading nginx ..."
docker-compose -f ./docker-compose-create-ssl.yml exec nginx nginx -s
reload

```

5. SSH into the Server

1. Use the following command to log in to your server:

```
ssh root@IP_ADDRESS
```

Replace **IP_ADDRESS** with the public IP address of your server.

This command connects you to your server as the root user.

Once logged in:

6. Set Up Git Configuration

Configure your Git to pull from your private repository using SSH:

1. **Generate an SSH key on your server:** Go to .ssh folder on the server:

```
cd .ssh
```

```
ssh-keygen -t rsa -b 4096 -C "OPTINAL COMMENT"
```

This command creates a 4096-bit RSA key pair with a comment containing your email address.

2. Copy the public key:

```
cat ~/.ssh/id_rsa.pub
```

Use this command to display the public portion of your SSH key.

3. Add the key to your GitHub account:

- Go to **GitHub > Settings > SSH and GPG keys**.
- Click **New SSH Key**, paste the public key, and save it.

Test the SSH connection:

```
ssh -T git@github.com
```

If successful, you should see a message like:

```
Hi username! You've successfully authenticated.
```

This confirms that your server can securely connect to GitHub via SSH.

7. Deploy the App and Create SSL

1. Add **init-letsencrypt.sh** to your **.gitignore** file:

- Open the **.gitignore** file in the root folder of your project in your local system.
- Add the following line:

```
init-letsencrypt.sh  
nginx/certbot
```

- This ensures that the sensitive files/folders are not pushed to the GitHub repository.

2. Push Your Changes to GitHub:

- Commit and push the updates:

```
git add .
git commit -m "Add SSL setup and ignore init-letsencrypt.sh"
git push origin main
```

3. Pull the Changes on the Server:

- First, ensure **Git** is installed on the server. If it's not installed, use the following command:

```
apt update
apt install -y git
```

- Navigate to the `/var/www/app` directory on your server:

```
cd /var/www/app
```

- Add your GitHub repository as the `origin` using SSH:

```
git init
git remote add origin git@github.com:yourusername/your-repo-
name.git
```

- Replace `yourusername/your-repo-name` with the actual username and repository name of your GitHub project.

- Pull the latest changes from your GitHub repository:

```
# master can be replaced with the actual branch name in your
repository
git pull origin master
```

4. Copy `init-letsencrypt.sh.sample` to `init-letsencrypt.sh`:

- Create the working script from the sample file and make it executable:

```
cp init-letsencrypt.sample.sh init-letsencrypt.sh
chmod +x init-letsencrypt.sh
```

- Update `APP_URL` and `EMAIL` in the copied file.

5. Install Docker and Docker Compose on the Server:

- If Docker and Docker Compose are not already installed, use the following commands to install them:

Install Docker:

```
apt update
apt install -y docker.io
systemctl start docker
systemctl enable docker
```

Install Docker Compose:

```
apt install -y curl
curl -L
"https://github.com/docker/compose/releases/latest/download/docke
r-compose-$(uname -s)-$(uname -m)" -o /usr/local/bin/docker-
compose
chmod +x /usr/local/bin/docker-compose
```

- Verify that Docker and Docker Compose are installed correctly:

```
docker --version
docker-compose --version
```

6. Run the `init-letsencrypt.sh` Script:

- Execute the script to create and configure SSL certificates:

```
./init-letsencrypt.sh
```

8. Verify SSL Creation

Once the script completes successfully, you will see an appropriate success message in the terminal, and you are ready to move to the next step.

8. Folder and File Structure

At the end of this step, your project should have the following folder and file structure:

```
.
├── nginx/                # Contains Nginx configurations and
└── Dockerfiles
```

```

|   |— default-create-ssl.conf      # Minimal Nginx configuration for
creating SSL certificates
|   |— default-dev.conf            # Nginx configuration for the
development environment
|   |— gzip.conf                  # Gzip compression configuration
|   |— Dockerfile.create-ssl      # Dockerfile for creating SSL
certificates
|   |— Dockerfile.dev             # Dockerfile for the Nginx service in
development
|— site/                          # Contains website files
|   |— index.html                 # Main HTML file
|   |— static/                   # Folder for static assets
|       |— css/
|           |— styles.css        # Test CSS file
|       |— js/
|           |— main.js          # Test JavaScript file
|       |— images/              # (Optional) Folder for image files
|— init-letsencrypt.sh.sample     # Sample SSL creation script
|— init-letsencrypt.sh           # Actual SSL creation script (added by
the user)
|— .gitignore                    # Git ignore file
|— docker-compose-create-ssl.yml  # Docker Compose file for creating SSL
certificates
|— docker-compose-dev.yml        # Docker Compose file for development
setup

```

Step 04: Running Application on The Server

Goal of This Step

In this step, we will:

1. Configure Nginx for a production environment with SSL.
2. Create the necessary Docker configurations to serve the app securely.
3. Set up a script to automate the deployment process.
4. Ensure the app is accessible through https://APP_URL.

By the end of this step, your app will run in production mode, fully secured with SSL.

Steps to Configure the Production Environment

1. Create the Nginx Configuration for Production

1. In the `nginx` folder, create a file named `default-prod-ssl.conf` with the following content:

```

# Redirects HTTP traffic to HTTPS
server {
    # Listens for HTTP traffic on port 80
    listen 80;

```

```
# Replace test4tutorial.work with your actual domain
server_name test4tutorial.work www.test4tutorial.work;

# Hides Nginx version info for security
server_tokens off;

# Handles Let's Encrypt requests for SSL verification
location /.well-known/acme-challenge/ {
    # Allows all traffic to this location
    allow all;

    # Serves the challenge files from this directory
    root /var/www/certbot;
}

# Redirects all traffic to HTTPS
location / {
    return 301 https://$host$request_uri;
}

# Redirects traffic from www.test4tutorial.work to test4tutorial.work over
# HTTPS
server {
    # Listens for HTTPS traffic on port 443
    listen 443 ssl;

    # Redirects requests to the non-www version of the domain
    server_name www.test4tutorial.work;

    # SSL certificate and key paths generated by Let's Encrypt
    ssl_certificate
/etc/letsencrypt/live/test4tutorial.work/fullchain.pem;
    ssl_certificate_key
/etc/letsencrypt/live/test4tutorial.work/privkey.pem;

    # Includes SSL configuration provided by Let's Encrypt
    include /etc/letsencrypt/options-ssl-nginx.conf;
    ssl_dhparam /etc/letsencrypt/ssl-dhparams.pem;

    # Redirects requests to the non-www domain
    return 301 https://test4tutorial.work$request_uri;
}

# Handles HTTPS traffic for test4tutorial.work
server {
    # Listens for HTTPS traffic on port 443
    listen 443 ssl;

    # Replace test4tutorial.work with your actual domain
    server_name test4tutorial.work;

    # SSL certificate and key paths generated by Let's Encrypt, and
```

```
Replace test4tutorial.work with your actual domain
    ssl_certificate
/etc/letsencrypt/live/test4tutorial.work/fullchain.pem;
    ssl_certificate_key
/etc/letsencrypt/live/test4tutorial.work/privkey.pem;

# Includes SSL configuration provided by Let's Encrypt
include /etc/letsencrypt/options-ssl-nginx.conf;
ssl_dhparam /etc/letsencrypt/ssl-dhparams.pem;

# Specifies the root directory for static and HTML files
root /var/www/app;
index index.html;

# Custom 404 error page
error_page 404 /404.html;

# Includes Gzip compression configuration
include /etc/nginx/conf.d/gzip.conf;

# Serves static files (CSS, JS, images, etc.) from the /static/ folder
location /static/ {
    alias /var/www/app/static;
}

# Configures long-term caching for images
location ~* \.(gif|jpe?g|png|svg)$ {
    expires max;
    add_header Cache-Control "public, max-age=31536000, immutable";
}

# Configures no-cache headers for CSS, JS, and JSON files
location ~* \.(css|js|json)$ {
    add_header Cache-Control "no-cache, must-revalidate";
}

# Serves requested files or falls back to a 404 error
location / {
    try_files $uri $uri.html $uri/ /404.html;
}
}
```

2. Create the Dockerfile for Production

1. In the `nginx` folder, create a file named `Dockerfile.prod-ssl` and add the following content:

```
# Use the lightweight Alpine-based Nginx image as the base image
FROM nginx:1.20.2-alpine

# Copy the production Nginx configuration into the container
```



```
COPY ./default-prod-ssl.conf /etc/nginx/conf.d/default.conf

# Copy the Gzip compression configuration
COPY ./gzip.conf /etc/nginx/conf.d/gzip.conf

# Create the static directory in the container
RUN mkdir -p /var/www/app/static && \
    chmod 755 /var/www/app/static
```

3. Create the Docker Compose File for Production

1. In the root folder of your project, create a file named `docker-compose-prod-ssl.yml` and add the following content:

```
services:
  nginx:
    # Restart the container automatically in case of failure
    restart: always

    # Build the Nginx image using the production Dockerfile
    build:
      context: ./nginx
      dockerfile: Dockerfile.prod-ssl

    # Map ports 80 (HTTP) and 443 (HTTPS) from the container to the
    server
    ports:
      - "80:80"
      - "443:443"

    # Mount necessary volumes
    volumes:
      - ./site:/var/www/app # Static files and HTML
      - ./nginx/certbot/www:/var/www/certbot # Let's Encrypt challenge
      directory
      - ./nginx/certbot/conf:/etc/letsencrypt # SSL certificates

  certbot:
    # Use the Certbot Docker image for SSL certificate management
    image: certbot/certbot:latest
    container_name: certbot

    # Share volumes for SSL certificates and challenges
    volumes:
      - ./nginx/certbot/www:/var/www/certbot
      - ./nginx/certbot/conf:/etc/letsencrypt

    # Runs Certbot in a loop to renew SSL certificates periodically
    entrypoint: "/bin/sh -c 'trap exit TERM; while ;; do sleep 6h &
    wait $$!}; certbot renew; done'"
```

4. Create the Deployment Script

1. In the root folder of your project, create a file named `run_app.sh` and make it executable:

```
chmod +x run_app.sh
```

2. Add the following content to `run_app.sh`:

```
# Pull the latest changes from the GitHub repository
git pull origin master

# Remove all running and stopped containers
docker container rm -f $(docker container ls -a -q)

# Remove all unused Docker images
docker image rm -f $(docker image ls -a -q)

# Stop and remove the current Docker Compose stack for SSL creation (if running)
docker-compose -f docker-compose-create-ssl.yml down

# Stop and remove the current Docker Compose stack for production (if running)
docker-compose -f docker-compose-prod-ssl.yml down

# Build and start the production Docker Compose stack
# Removes unused Docker volumes.
# Removes all unused build cache data.
# Removes all dangling and unused Docker images.
# Removes all stopped containers.
# Removes all unused Docker networks.
docker-compose -f docker-compose-prod-ssl.yml up --build -d && docker
volume prune -f && docker builder prune -a -f && docker image prune -a -f
&& docker container prune -f && docker network prune -f
```

5. Push the Changes to GitHub

1. Add all the newly created files to the Git repository:

```
git add .
git commit -m "Add production setup with SSL configuration"
git push origin main
```

6. Deploy the App on the Server

1. Pull the Latest Changes:

- Log into your server and navigate to the `/var/www/app` directory:

```
cd /var/www/app
```

- Pull the latest changes from the GitHub repository:

```
git pull origin master
```

2. Run the Deployment Script:

- Execute the deployment script to set up and run the production environment:

```
./run_app.sh
```

7. Verify Your App

1. Check App Accessibility:

- Visit `https://APP_URL` (replace `APP_URL` with your actual domain name).
- Verify that your app is accessible and loading over HTTPS.

2. Confirm SSL Configuration:

- Ensure the connection is secure, indicated by a padlock icon in the browser's address bar.
- Both `https://APP_URL` and `https://www.APP_URL` should be functional.

3. Verify Static Assets:

- Ensure that all static files (e.g., CSS, JavaScript, images) are being served correctly.
- Check browser developer tools (**Network tab**) to confirm files are served with appropriate caching and compression.

8. Automate Nginx Reload for Certificate Renewal

When Certbot renews SSL certificates, Nginx needs to be reloaded to apply the updated certificates. To automate this process, you can add a cron job to reload Nginx every 24 hours.

1. Open the Crontab:

```
crontab -e
```

2. Add the Following Cron Job:

```
0 0 * * * docker exec app-nginx-1 nginx -s reload
```

- This command will reload Nginx at midnight every day.
- Replace app-nginx-1 with the name of your Nginx container if it's different.

3. Save and Exit:

- Save the changes to your crontab file and exit the editor.

4. Verify the Cron Job:

- To check if the cron job has been added, use:

```
crontab -l
```

- This will list all the active cron jobs for the current user.

With this cron job, Nginx will automatically reload every 24 hours, ensuring that renewed SSL certificates are applied without manual intervention.

9. Folder and File Structure

At the end of this step, your project folder structure should look like this:

```
.
├── nginx/                                # Contains Nginx configurations and
Dockerfiles                               #
├── ├── default-prod-ssl.conf              # Nginx configuration for production
with SSL                                  #
├── ├── default-create-ssl.conf            # Nginx configuration for creating SSL
certificates                             #
├── ├── default-dev.conf                   # Nginx configuration for the
development environment                  #
├── ├── gzip.conf                          # Gzip compression configuration
├── ├── Dockerfile.prod-ssl                # Dockerfile for the Nginx production
setup                                    #
├── ├── Dockerfile.create-ssl              # Dockerfile for SSL creation setup
├── └── Dockerfile.dev                     # Dockerfile for the Nginx development
setup                                    #
├── site/                                 # Contains website files
└── └── index.html                         # Main HTML file
```

```
├── static/                                # Folder for static assets
│   ├── css/
│   │   └── styles.css                    # Test CSS file
│   ├── js/
│   │   └── main.js                      # Test JavaScript file
│   └── images/                          # (Optional) Folder for image files
├── init-letsencrypt.sh.sample            # Sample SSL creation script
├── init-letsencrypt.sh                   # Actual SSL creation script (added by
the user)
├── run_app.sh                           # Deployment script for the production
environment
├── .gitignore                           # Git ignore file
├── docker-compose-prod-ssl.yml           # Docker Compose file for production
setup
├── docker-compose-create-ssl.yml         # Docker Compose file for creating SSL
certificates
└── docker-compose-dev.yml                # Docker Compose file for development
setup
```

Conclusion

Throughout this tutorial, we have covered the foundational steps to set up a functional and efficient CI/CD pipeline for a static website. Here's a summary of what we accomplished:

- **Step 01:** We configured the **development environment**, learning how to set up a basic structure with Nginx and Docker to serve a static website locally.
- **Step 02:** We enhanced the setup by adding support for **static files** (e.g., CSS, JavaScript, images) and enabled **Gzip compression** to improve performance.
- **Step 03:** We built a setup for creating **SSL certificates** using Certbot, ensuring that the website could be served securely over HTTPS.
- **Step 04:** We completed the **production environment** setup, creating a robust and scalable solution for hosting our static website with Nginx, Docker, and SSL.

This tutorial has been designed as a **step-by-step guide**, focusing on **implementation-first learning** to help you understand how to configure different setups for development, SSL generation, and production environments. Rather than diving deep into every command of Nginx or Docker, the emphasis was on applying key commands in real-world scenarios, allowing you to digest the concepts more effectively.

The Bigger Picture

This tutorial serves as a **foundational step** for building more sophisticated CI/CD pipelines. While we implemented it on a simple static website, the approach and structure introduced here can be applied to more complex systems. Future tutorials can build upon this base to include:

- **Full-stack development setups** that integrate:
 - Front-end frameworks like React or Angular.

- Back-end systems like Django or Node.js.
 - Task queues like Celery, and additional services like Redis and PostgreSQL.
 - Advanced CI/CD pipelines with automated testing, staging environments, and deployment workflows.
-

Next Steps and Areas for Improvement

To further enhance this structure, here are some advanced areas you can explore:

1. Implement Docker Swarm

- Docker Swarm is more efficient than Docker Compose for complex structures because it enables **distribution across multiple nodes** and provides **near zero-downtime deployment**.
 - For this simple static website, Docker Compose works perfectly fine. However, for more complex systems, Docker Swarm (or Kubernetes) becomes essential.
 - **Assignment:** Spend some time learning about Docker Swarm and feel free to reach out if you have questions. I'll cover it in future courses for more advanced systems.
-

2. Automate Further

- You can create **custom shell scripts** that act as agents, automating the deployment process even further. These scripts can:
 - Log in to the server and pull changes from GitHub automatically.
 - Manage versioning and update the production environment without requiring manual intervention.
 - **Assignment:** Research and experiment with automated deployment agents. This will help you build more scalable and efficient deployment processes.
-

3. Server Hardening

- Security is critical in a production environment. Some best practices include:
 - **Disable root login:** Use SSH keys and create a specific user for added security.
 - **Enable firewalls:** Use tools like **ufw** to manage access to your server.
 - **Apply server hardening techniques** to protect against unauthorized access.
 - **Assignment:** Learn about server hardening and apply these best practices to your setup.
-

4. Build a More Complex System

- In this tutorial, we focused on a static website. The next step is to implement a more complex system, such as:
 - Combining **front-end and back-end technologies**.
 - Using tools like React, Django, Celery, and Redis to simulate real-world, large-scale applications.
 - **Assignment:** Start exploring full-stack systems. If you'd like me to create a course for such setups, let me know!
-

Your Feedback Matters

Your support and feedback are invaluable. They inspire me to create more detailed tutorials like this, focusing on progressively mastering more sophisticated structures. I'd love to hear your thoughts about this course:

- Was the step-by-step approach helpful?
- What other topics or systems would you like to learn about?

Feel free to share your feedback and suggestions. Together, we can continue building a strong foundation and advance to mastering complex development and deployment setups!

Thank you for following along, and I hope this tutorial helps you build your skills and confidence in configuring and deploying web applications!